

5. DriveTrain in MCP

DriveTrain in MCP

Chapter Details	
Chapter Goal	Deploy and use a Drivetrain cluster
Chapter Sections	6.1. Prepare virtual lab 5.2. Deploy Drivetrain 5.3. Use Drivetrain

5.1. Prepare virtual lab

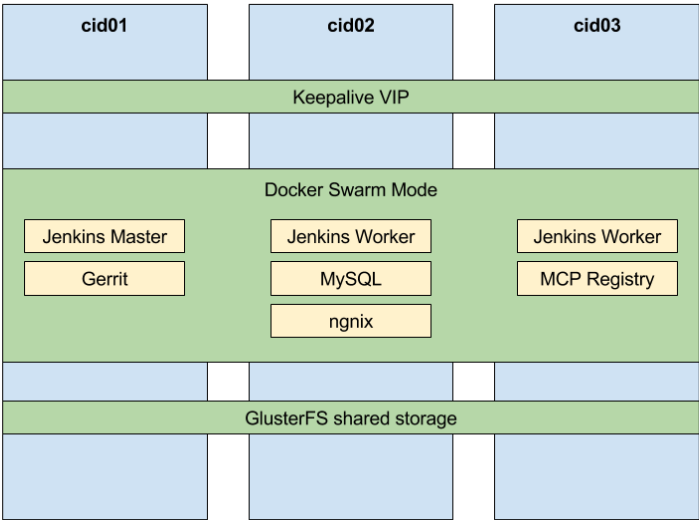
5.1.1. Environment Overview

MCP Drivetrain consists of:

- Jenkins
- Gerrit
- Aptly

Those are integrated together with reclass model.

MCP deployment guide recommends that DriveTrain deployment is done on 3 virtual nodes, where each service (Jenkins, Aptly and Gerrit) runs inside Docker container and those containers are managed by Docker swarm mode. Each of those components uses GlusterFS as a distributed storage for storing service-specific data reliably in distributed fashion. Diagram below shows the recommended deployment model below:



5.1.2. Configure networking

Step 1 Download the model and scripts:

```
stack@lab:~$ git clone --recurse-submodules git@bitbucket.org:mirantis-training/mcp100-
```

Step 2 Install prerequisites:

```
stack@lab:~$ sudo apt-get update
```

```
stack@lab:~$ sudo apt-get install -y git mkisofs curl virtinst cpu-checker qemu-kvm
```

Step 3 Create virtual networks:

```
stack@lab:~$ mcp100-drivetrain/provision-networks.sh
```

5.1.3. Prepare cfg01

Step 1 Create a model:

```
stack@lab:~$ cd mcp100-drivetrain/
```

```
stack@lab:~/mcp100-drivetrain$ sudo su
```

```
root@lab:/home/stack/mcp100-drivetrain# mkdir /root/model
```

```
root@lab:/home/stack/mcp100-drivetrain# cp -rT /home/stack/mcp100-drivetrain/ /root/moc
```

Step 2 Download images:

```
root@lab:/home/stack/mcp100-drivetrain# wget http://repos/cfg01-day01-2018.8.0.qcow2 -C  
/pool/images/cfg01.drivetrain-ha.local.img
```

Step 3 Prepare directory for configuration ISO. Set MCP version and clone pipeline repositories:

```
root@lab:/home/stack/mcp100-drivetrain# export version="2018.11.0"
```

```
root@lab:/home/stack/mcp100-drivetrain# git clone https://github.com/Mirantis/mk-pipeli
```

```
Cloning into '/root/mk-pipelines'...  
remote: Counting objects: 4064, done.  
remote: Compressing objects: 100% (86/86), done.  
remote: Total 4064 (delta 66), reused 86 (delta 36), pack-reused 3942  
Receiving objects: 100% (4064/4064), 1.02 MiB | 0 bytes/s, done.  
Resolving deltas: 100% (2612/2612), done.  
Checking connectivity... done.
```

```
root@lab:/home/stack/mcp100-drivetrain# git clone https://github.com/Mirantis/pipeline-
```

```
Cloning into '/root/pipeline-library'...  
remote: Counting objects: 6401, done.  
remote: Compressing objects: 100% (40/40), done.  
remote: Total 6401 (delta 26), reused 63 (delta 16), pack-reused 6323  
Receiving objects: 100% (6401/6401), 927.30 KiB | 0 bytes/s, done.  
Resolving deltas: 100% (2924/2924), done.  
Checking connectivity... done.
```

Step 4 Checkout specific version of pipelines:

```
root@lab:/home/stack/mcp100-drivetrain# pushd /root/mk-pipelines && git checkout tags/5.0.0
~/mk-pipelines ~ /home/stack/mcp100-drivetrain
```

Note: checking out 'tags/2018.3.1'.

You are in 'detached HEAD' state. You can look around, make experimental changes and commit them, and you can discard any commits you make in this state without impacting any branches by performing another checkout.

If you want to create a new branch to retain commits you create, you may do so (now or later) by using `-b` with the checkout command again. Example:

```
git checkout -b <new-branch-name>
```

```
HEAD is now at 5dac8d8... cloud-update pipeline rename Merges param
/home/stack
```

```
root@lab:/home/stack/mcp100-drivetrain# pushd /root/pipeline-library && git checkout tags/2018.3.1
~/pipeline-library /home/stack
```

Note: checking out 'tags/2018.3.1'.

You are in 'detached HEAD' state. You can look around, make experimental changes and commit them, and you can discard any commits you make in this state without impacting any branches by performing another checkout.

If you want to create a new branch to retain commits you create, you may do so (now or later) by using `-b` with the checkout command again. Example:

```
git checkout -b <new-branch-name>
```

```
HEAD is now at 658b110... Merge "Change default version of salt-models-testing docker image"
/home/stack
```

Step 5 Download script *create_config_drive.sh*:

```
root@lab:/home/stack/mcp100-drivetrain# wget -O /root/create-config-drive \
https://raw.githubusercontent.com/Mirantis/mcp-common-scripts/${version}/config-drive/c
```

```
root@lab:/home/stack/mcp100-drivetrain# chmod +x /root/create-config-drive
```

The script will be used to create configuration ISO

Step 6 Download cloud-init script for **cfg01**:

```
root@lab:/home/stack/mcp100-drivetrain# wget -O /root/user_data.sh \
https://raw.githubusercontent.com/Mirantis/mcp-common-scripts/${version}/config-drive/n
```

Step 7 Edit file `/root/user_data.sh`, modify environment parameters as shown below:

```
write_files:
- owner: root:root
  path: /etc/cloud/master_environment
  permissions: '0644'
  content: |
    [ -f /etc/cloud/master_environment_override ] && . /etc/cloud/master_environment_
    export SALT_MASTER_DEPLOY_IP=192.168.10.100
    export SALT_MASTER_MINION_ID=cfg01.drivetrain-ha.local
    export DEPLOY_NETWORK_GW=192.168.10.1
    export DEPLOY_NETWORK_NETMASK=255.255.255.0
    export DNS_SERVERS=172.19.0.6
```

```
export DEPLOY_NETWORK_MTU=1500
export MCP_VERSION=2018.11.0
```

Step 8 Create configuration ISO using downloaded script from previous steps:

```
root@lab:/home/stack/mcp100-drivetrain# /root/create-config-drive -u /root/user_data.sh
--model /root/model --mk-pipelines /root/mk-pipelines \
--pipeline-library /root/pipeline-library /pool/images/cfg01.drivetrain-ha.local.iso

adding user data from /root/user_data.sh
adding reclass model directory /root/model
adding mk_pipelines directory /root/mk-pipelines
adding pipeline_library directory /root/pipeline-library
generating configuration image at /var/lib/libvirt/images/cfg01/cfg01-config.local.iso
```

5.1.4. Create cfg01 VM

Step 1 Create VM:

```
root@lab:/home/stack/mcp100-drivetrain# virt-install --name cfg01.drivetrain-ha.local \
--disk path=/pool/images/cfg01.drivetrain-ha.local.img,bus=virtio,format=qcow2,cache=nc
--disk path=/pool/images/cfg01.drivetrain-ha.local.iso,device=cdrom \
--network network:br_pxe,model=virtio \
--network network:br_mgmt,model=virtio \
--network network:br_ctl,model=virtio \
--network network:br_ext,model=virtio \
--ram 8192 --vcpus=4 --accelerate \
--boot hd --vnc --noreboot --autostart

Starting install...
Creating domain...
Failed to connect to Mir: Failed to connect to server socket: No such file or directory
Unable to init server: Could not connect: Connection refused
Cannot open display:
Run 'virt-viewer --help' to see a full list of available command line options
Domain creation completed. You can restart your domain by running:
  virsh --connect qemu:///system start cfg01.drivetrain-ha.local
```

Step 2 Update networking:

```
root@lab:/home/stack/mcp100-drivetrain# virsh net-update br_pxe add ip-dhcp-host "\
<host mac='${$(virsh domiflist cfg01.drivetrain-ha.local | grep br_pxe | awk '{print $5}'}'
" --live --config

Updated network br_pxe persistent config and live state
```

Step 3 Start cfg01 VM:

```
root@lab:/home/stack/mcp100-drivetrain# virsh start cfg01.drivetrain-ha.local

Domain cfg01.drivetrain-ha.local started
```

Step 4 Enter *cfg01* and wait until cloud-init is done. It should take around 7 min to complete:

```
root@lab:/home/stack/mcp100-drivetrain# virsh console cfg01.drivetrain-ha.local

Connected to domain cfg01.drivetrain-ha.local
Escape character is ^]
```

```
[ OK ] Started LSB: Machine Check Exceptions (MCE) collector & decoder.
[ OK ] Started Login Service.
      Starting Daily apt upgrade and clean activities...
      Starting Terminate Plymouth Boot Screen...
      Starting Hold until boot process finishes up...
[ OK ] Started LSB: Start NTP daemon.
```

Step 5 Update SSH keys:

```
root@cfg01:~# salt-call state.sls openssh
```

Exit to host lab using *CTRL+]*

5.1.5. Create other VMs

Step 1 Spawn cluster VMs:

```
root@lab:/home/stack/mcp100-drivetrain# ./provision-cluster.sh
```

```
root@lab:/home/stack/mcp100-drivetrain# virsh list
```

Id	Name	State
14	cfg01.drivetrain-ha.local	running
15	cid01.drivetrain-ha.local	running
16	cid02.drivetrain-ha.local	running
17	cid03.drivetrain-ha.local	running

Switch back to *stack* user:

```
root@lab:/home/stack/mcp100-drivetrain# exit
```

```
stack@lab:~/mcp100-drivetrain$
```

5.2. Deploy Drivetrain

The deployment of Drivetrain will require several salt states to be applied to cluster nodes.

5.2.1. Prepare Salt

Step 1 Login to *cfg01* using SSH:

```
stack@lab:~/mcp100-drivetrain$ ssh root@192.168.10.100
```

```
root@cfg01:~#
```

Step 2 Bootstrap cluster nodes:

```
root@cfg01:~# salt '*' saltutil.sync_all
```

```
root@cfg01:~# salt '*' state.sls linux.system.repo
```

```
root@cfg01:~# salt '*' state.sls salt.minion,linux,openssh,ntp
```

Step 3 Install GlusterFS:

```
root@cfg01:~# salt -C 'I@glusterfs:server' state.sls glusterfs.server.service
root@cfg01:~# salt -C 'I@glusterfs:server:role:primary' state.sls glusterfs.server.setu
```

Verify GlusterFS server:

```
root@cfg01:~# salt -C 'I@glusterfs:server' cmd.run 'gluster peer status; gluster volume
```

Ensure GlusterFS clusters is ready:

```
root@cfg01:~# salt -C 'I@glusterfs:client' state.sls glusterfs.client
```

Step 4 Install HAProxy:

```
root@cfg01:~# salt -C 'I@haproxy:proxy and I@docker:host' state.sls haproxy,keepalived
```

Step 5 Install Docker:

```
root@cfg01:~# salt -C 'I@docker:host' state.sls docker.host
```

Step 6 Enable swarm master and update salt mines:

```
root@cfg01:~# salt -C 'I@docker:swarm:role:master' state.sls docker.swarm
root@cfg01:~# salt -C 'I@docker:swarm' state.sls salt
root@cfg01:~# salt -C 'I@docker:swarm' mine.flush
root@cfg01:~# salt -C 'I@docker:swarm' mine.update
root@cfg01:~# salt -C 'I@docker:swarm' saltutil.sync_all
```

Step 7 Enable swarm workers:

```
root@cfg01:~# salt -C 'I@docker:swarm' state.sls docker.swarm
```

Update mines:

```
root@cfg01:~# salt -C 'I@docker:swarm' mine.update
```

Join swarm nodes:

```
root@cfg01:~# salt -C 'I@docker:swarm' state.sls docker.swarm
```

Verify swarm:

```
root@cfg01:~# salt -C 'I@docker:swarm:role:master' cmd.run 'docker node ls'
```

Step 8 Install Aptly:

```
root@cfg01:~# salt -C 'I@aptly:publisher' state.sls aptly.publisher
root@cfg01:~# salt -C 'I@docker:swarm:role:master' state.sls docker.client
```

Check that all replicas have been spawned:

```
root@cfg01:~# salt -C 'I@docker:swarm:role:master' cmd.run 'docker service ls'
```

Step 9 Update pillars:

```
root@cfg01:~# salt '*' saltutil.sync_all
```

Step 10 Install openldap:

```
root@cfg01:~# salt -C 'I@openldap:client' state.sls openldap
```

Step 11 Install gerrit:

```
root@cfg01:~# salt -C 'I@gerrit:client' state.sls gerrit
```

Step 12 Install jenkins:

```
root@cfg01:~# salt -C 'I@jenkins:client' state.sls jenkins
```

5.3. Use Drivetrain

5.3.1. Create ad-hoc formula with a static file

Step 1 Create dir:

```
root@cfg01:~# mkdir /srv/salt/env/prd/hello
```

Step 2 Create state /srv/salt/env/prd/hello/init.sls file:

```
/etc/hello.conf:
file.managed:
- source: salt://hello/files/hello.conf
- user: root
- group: root
- mode: 644
```

Step 3 Create dir:

```
root@cfg01:~# mkdir /srv/salt/env/prd/hello/files
```

Step 4 Create a static file /srv/salt/env/prd/hello/files/hello.conf:

```
This is hello file version 1
```

Step 5 Run state:

```
root@cfg01:~# salt "cfg01*" state.sls hello
```

Step 6 Verify the result:

```
root@cfg01:~# cat /etc/hello.conf
```

```
This is hello file version 1
```

5.3.2. Store the formula on git and install it from git using your cluster model

Step 1 Create git repo directory:

```
root@cfg01:~# mkdir /root/hello_repo
```

```
root@cfg01:~# cd /root/hello_repo
```

Step 2 Create empty git repo:

```
root@cfg01:~# git init
```

```
root@cfg01:~# git config --global user.email "student@mirantis.com"
```

Step 3 Copy formula:

```
root@cfg01:~# cp -r /srv/salt/env/prd/hello/ .
```

Step 4 Increment version in the config file hello/files/hello.conf:

```
This is hello file version 2
```

Step 5 Add and commit formula files to the git repo:

```
root@cfg01:~# git add -A
```

```
root@cfg01:~# git commit -m 'Initial commit'
```

Step 6 Modify your cluster model to install the formula from git repo. Edit /srv/salt/reclass/classes/cluster/drivetrain-ha/infra/hello.yml:

```
parameters:
  salt:
    master:
      environment:
        prd:
          formula:
            hello:
              source: git
              name: salt-formula-hello
              address: /root/hello_repo
              branch: master
```

Edit /srv/salt/reclass/classes/cluster/drivetrain-ha/infra/config/init.yml:

```
classes:
- cluster.drivetrain-ha.infra.hello
```

Step 7 Sync pillars:

```
root@cfg01:~# salt \* saltutil.sync_all
```

Step 8 Check if salt state installs your hello formula:

```
root@cfg01:~# salt "cfg01*" state.show_sls salt --output=yaml | less
```


Step 9 Install your hello fomula:

```
root@cfg01:~# salt "cfg01*" state.sls salt
```

Step 10 Check if the old ad-hoc formula is replaced with the new one from git repo:

```
root@cfg01:~# ls -l /srv/salt/env/prd/hello

lrwxrwxrwx 1 root root 50 Sep  2 13:52 /srv/salt/env/prd/hello -> /usr/share/salt-formu
root@cfg01:~# cat /srv/salt/env/prd/hello/files/hello.conf

This is hello file version 2
```

Step 11 Run your hello formula:

```
root@cfg01:~# salt "cfg01*" state.sls hello
```

Step 12 Verify the result:

```
root@cfg01:~# cat /etc/hello.conf

This is hello file version 2
```

5.3.3. Modify the formula, so hello.conf is parametrized with pillar

Step	1	Add	hello	version	parameter	(key-value	pair).	Edit
/srv/salt/reclass/classes/cluster/drivetrain-ha/infra/hello.yml:								

```
parameters:
  hello:
    version: 3
  salt:
    master:
      environment:
        prd:
          formula:
            hello:
              source: git
              name: salt-formula-hello
              address: /root/hello_repo
              branch: master
```

Step 2 Modify the formula:

```
root@cfg01:~# cd /root/hello_repo
```

Edit hello/init.sls:

```
/etc/hello.conf:
file.managed:
- source: salt://hello/files/hello.conf
- user: root
- group: root
- mode: 644
- template: jinja
```

Edit `hello/files/hello.conf`

```
This is hello file version {{ salt['pillar.get']('hello:version') }}
```

Step 3 Add and commit to git repo:

```
root@cfg01:~# git add -A
root@cfg01:~# git commit -m 'Hello ver from pillar'
```

Step 4 Sync:

```
root@cfg01:~# salt '*' saltutil.sync_all
```

Step 5 Install the new formula:

```
root@cfg01:~# salt "cfg01*" state.sls salt
```

Step 6 Run the formula:

```
root@cfg01:~# salt "cfg01*" state.sls hello
```

Step 7 Verify the result:

```
root@cfg01:~# cat /etc/hello.conf
```

```
This is hello file version 3
```

5.3.4. Store cluster model in a git repository and source it from there

Step 1 Create git repository:

```
root@cfg01:~# mkdir /root/model_repo
root@cfg01:~# cd /root/model_repo
root@cfg01:~# git init
root@cfg01:~# git config receive.denyCurrentBranch ignore
```

Step 2 Modify the model to use your git repo:

```
root@cfg01:~# cd /srv/salt/reclass
```

Edit `/srv/salt/reclass/classes/cluster/drivetrain-ha/infra/config/nodes.yml`:

```
parameters:
  reclass:
    storage:
      data_source:
        # don't use Git storage for model, instead use whatever is provided
        # with the lab
        # engine: local
        engine: git
        address: /root/model_repo
        branch: master
        force_reset: true
```

Step 3 Add and commit to the git repo in /srv/salt/reclass:

```
root@cfg01:~# git add -A
root@cfg01:~# git commit -m 'Switch storage to Git'
```

Reclass model note:

/srv/salt/reclass repo is created during bootstrapping MCP

Step 5 Push the model to your new model repo:

```
root@cfg01:~# git remote set-url origin /root/model_repo/
root@cfg01:~# git push -u origin master
```

Step 6 Regenerate the model in /srv/salt/reclass using your new model repo:

```
root@cfg01:~# salt-call state.sls reclass
```

Step 7 Modify the metadata used by your hello formula:

```
root@cfg01:~# cd /root/model_repo
root@cfg01:~# git checkout
```

Edit classes/cluster/drivetrain-ha/infra/hello.yml:

```
parameters:
  hello:
    version: 4
  salt:
    master:
      environment:
        prd:
          formula:
            hello:
              source: git
              name: salt-formula-hello
              address: /root/hello_repo
              branch: master
```

Step 7 Add and commit changes:

```
root@cfg01:~# git config --global user.email "student@mirantis.com"
root@cfg01:~# git add -A
root@cfg01:~# git commit -m 'Hello ver 4'
```

Step 8 Run your hello formula:

```
root@cfg01:~# salt "cfg01*" state.sls hello
```

Step 9 Verify the result:

```
root@cfg01:~# cat /etc/hello.conf
```

```
This is hello file version 3
```

Step 10 It should be still older version, we need to regenerate the model in /srv/salt/reclass:

```
root@cfg01:~# salt-call state.sls reclass
```

Step 11 Verify if the model is regenerated:

```
root@cfg01:~# cat /srv/salt/reclass/classes/cluster/drivetrain-ha/infra/hello.yml

parameters:
  hello:
    version: 4
```

Step 12 Run your hello formula:

```
root@cfg01:~# salt "cfg01*" state.sls hello
```

Step 13 Verify the result:

```
root@cfg01:~# cat /etc/hello.conf
```

```
This is hello file version 4
```

5.3.5. Use Jenkins GUI to run your hello formula

Step 1 Modify the formula metadata in the cluster model:

```
root@cfg01:~# cd /root/model_repo
```

Edit /root/model_repo/classes/cluster/drivetrain-ha/infra/hello.yml:

```
parameters:
  hello:
    version: 5
```

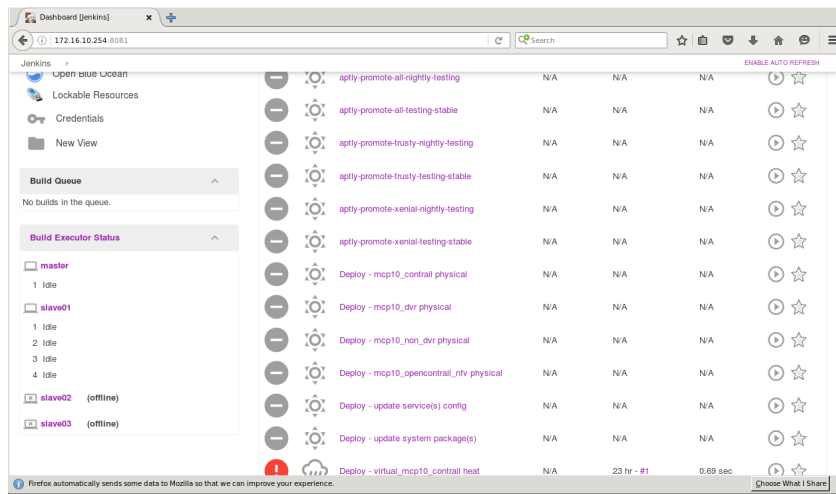
Step 2 Add and commit:

```
root@cfg01:~# git add -A
```

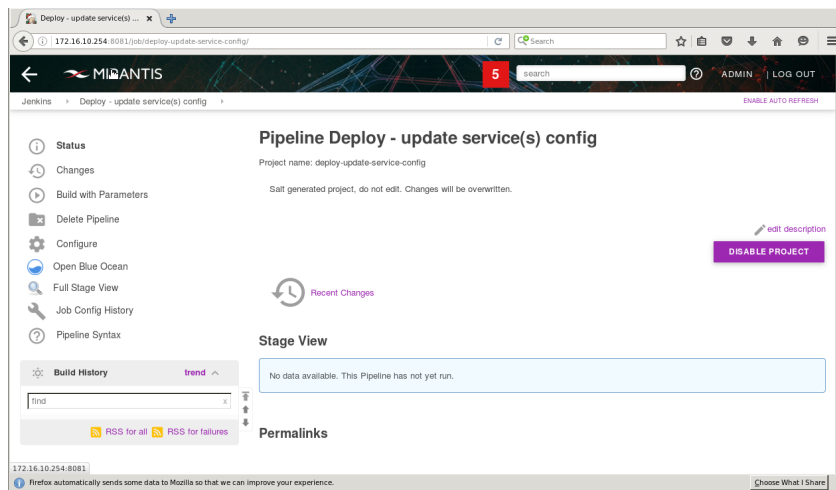
```
root@cfg01:~# git commit -m 'Hello ver 5'
```

Step 3 Open browser and navigate to <http://172.16.10.104:8081>

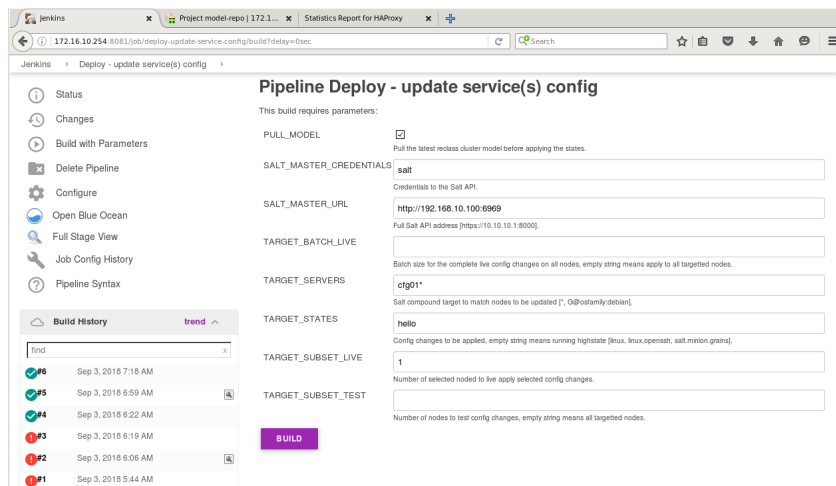
Use credentials **admin/password** to login and open list of pipelines:



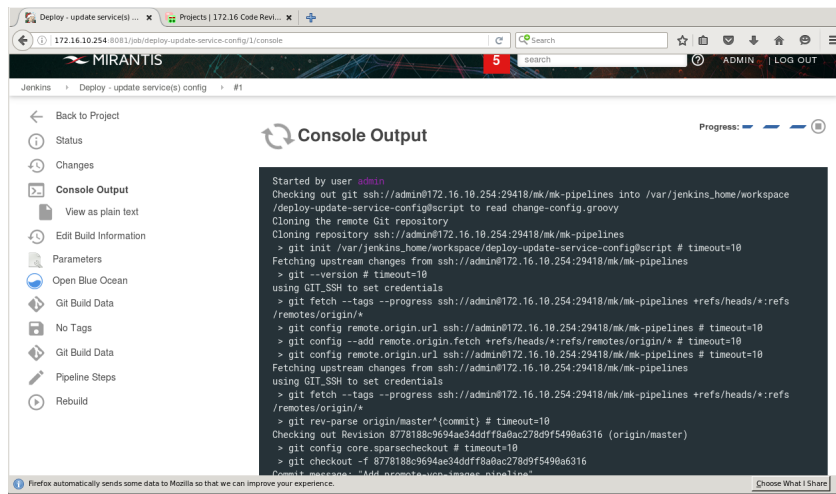
Find Deploy - Update service(s) config job:



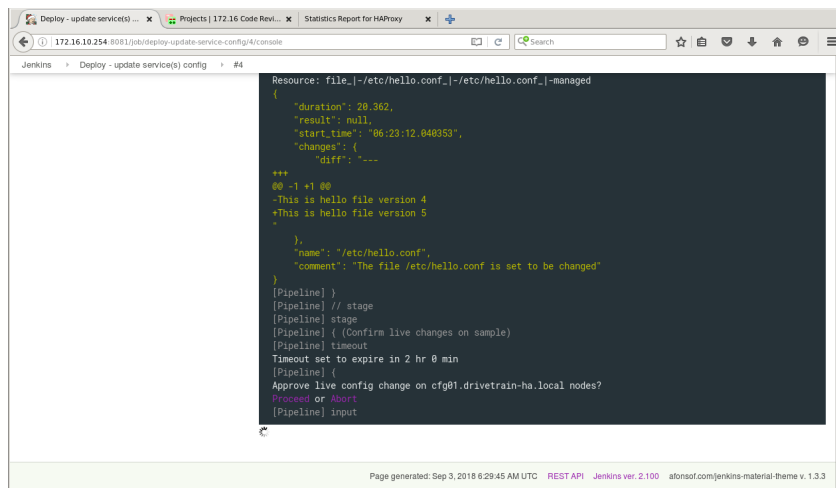
Launch the job with following parameters:



Open a console output to monitor pipeline execution:



When Pipeline asks to Approve config change click on **Approve**:



Step 4 Verify the result:

```
root@cfg01:~# cat /etc/hello.conf
```

```
This is hello file version 5
```

5.3.6. Use Jenkins API to run your hello formula

Step 1 Modify the metadata used by your hello formula:

```
root@cfg01:~# cd /root/model_repo
```

Edit `/root/model_repo/classes/cluster/drivetrain-ha/infra/hello.yml`

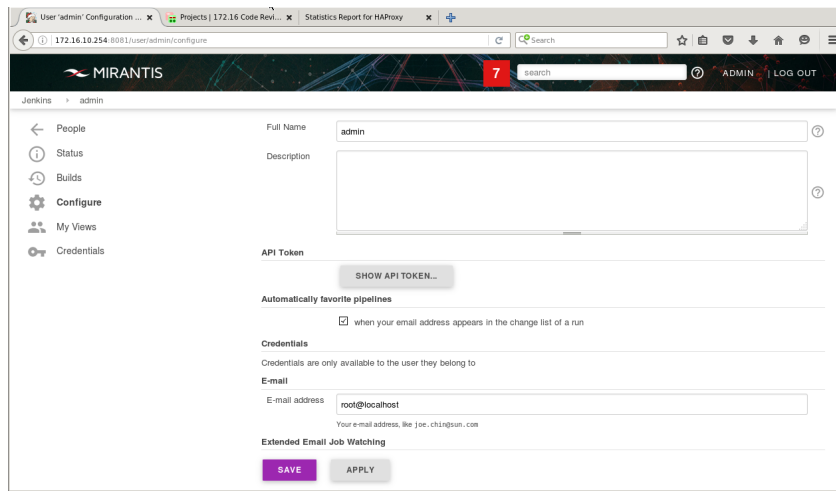
```
parameters:
  hello:
    version: 6
```

Step 2 Add and commit

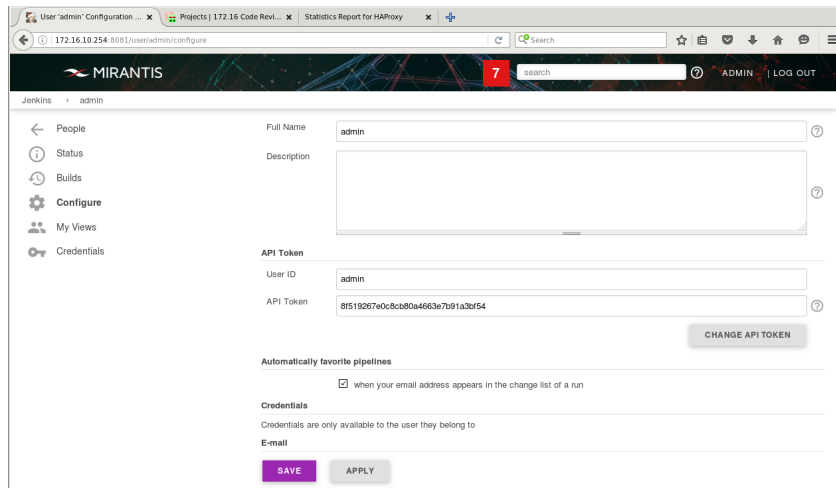
```
root@cfg01:~# git add -A
```

```
root@cfg01:~# git commit -m 'Hello ver 6'
```

Step 3 Get the user token. Click on Admin button at the right-top. :

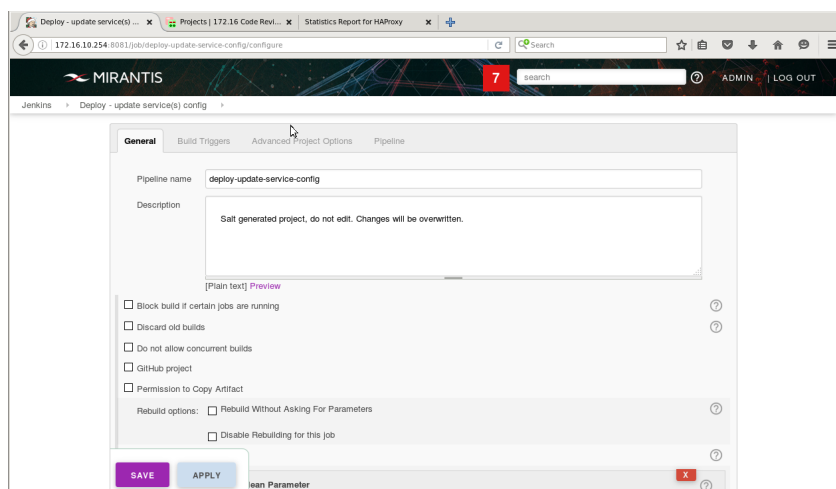


Open *Configure* tab and look for API token:



In our case it's admin/8f519267e0c8cb80a4663e7b91a3bf54

Step 4 Get back to the list of pipelines and open Configuration tab of pipeline “Deploy - Update service(s) config” and get the project/pipeline name:



It should be *deploy-update-service-config*

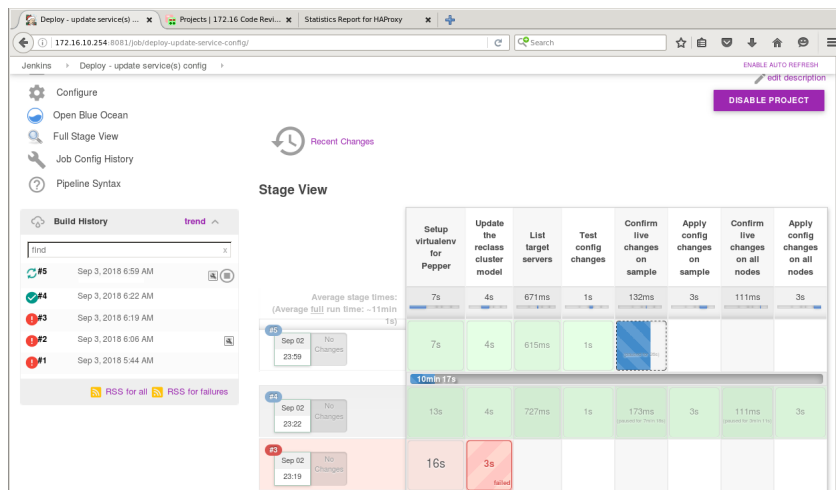
Step 5 Call the Jenkins API using userID/token/pipeline values:

```
root@cfg01:~# curl -d "TARGET_SERVERS=cfg01.drivetrain-ha.local" -d "TARGET_STATES=hell" -i -X POST http://admin:8f519267e0c8cb80a4663e7b91a3bf54@172.16.10.104:8081/job/deploy-
```

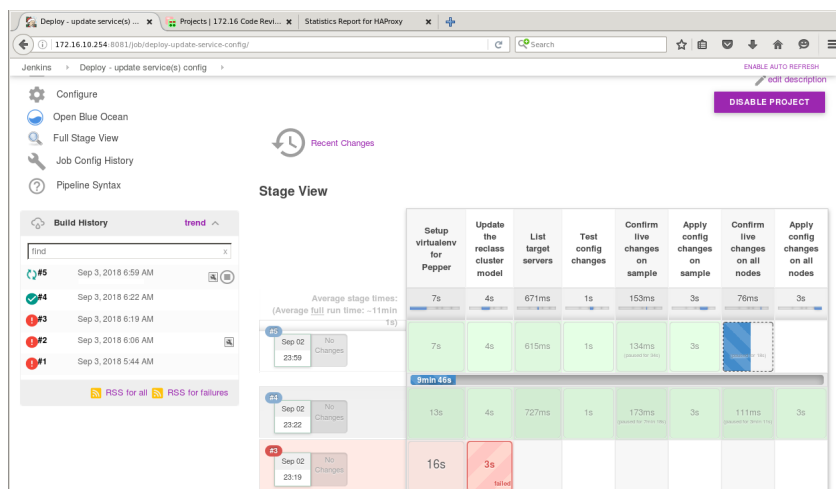
Remember this command - we will use it in next steps.

Step 6 Accept steps in the pipeline. This time you can do it on Stage View by clicking on Confirm changes stage:

Confirm live changes on sample:



Confirm live changes on all nodes:



Step 7 Verify the result:

```
root@cfg01:~# cat /etc/hello.conf
```

```
This is hello file version 6
```

5.3.7. Modify your git repo to automatically call the Jenkins pipeline

Step 1 Add a post commit hook:

```
root@cfg01:~# cd /root/model_repo
```

Edit file `/root/model_repo/.git/hooks/post-commit`

```
#!/bin/bash
curl -d "TARGET_SERVERS=cfg01.drivetrain-ha.local" -d "TARGET_STATES=hello" -i -X POST
```



```
root@cfg01:~# chmod 755 .git/hooks/post-commit
```

Step 2 Modify the metadata used by your hello formula. Edit file `/root/model_repo/classes/cluster/drivetrain-ha/infra/hello.yml`

```
parameters:
  hello:
    version: 7
```

Step 3 Add and commit:

```
root@cfg01:~# git add -A
root@cfg01:~# git commit -m 'Hello ver 7'
```

Step 4 Accept steps in the pipeline

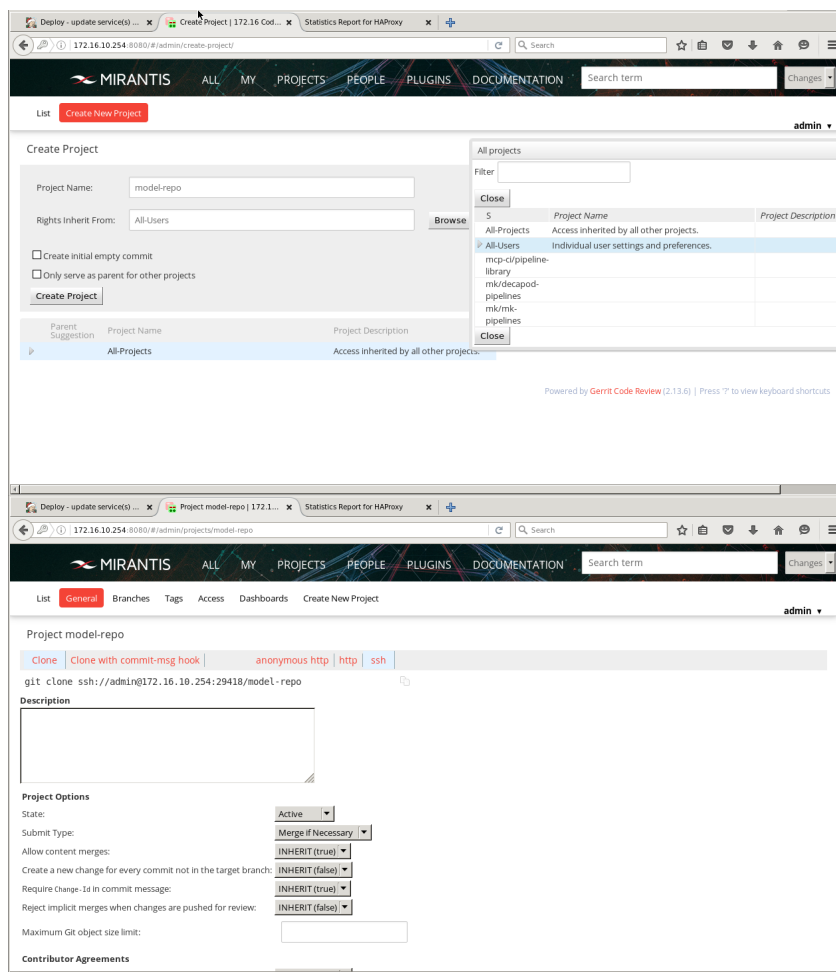
Step 5 Verify the results:

```
root@cfg01:~# cat /etc/hello.conf

This is hello file version 7
```

5.3.8. Move to a shared repo in Gerrit

Step 1 Open browser and navigate to `http://172.16.10.104:8080`. Enter **admin/password** as credentials. Create project `model-repo` in Gerrit



Step 2 Push to Gerrit repo:

```
root@cfg01:~# cd /root/model_repo/

root@cfg01:~# git remote add origin ssh://admin@172.16.10.104:29418/model-repo

root@cfg01:~# git push origin master

root@cfg01:~# rm .git/hooks/post-commit
```

Step 3 Modify your model to use Gerrit:

```
root@cfg01:~# cd /root/model_repo/
```

Edit `/srv/salt/reclass/classes/cluster/drivetrain-ha/infra/config/nodes.yml`:

```
parameters:
  reclass:
    storage:
      data_source:
        # don't use Git storage for model, instead use whatever is provided
        # with the lab
        # engine: local
        engine: git
        address: ssh://admin@172.16.10.104:29418/model-repo
        branch: master
        force_reset: true
```

Step 4 Add, commit and push to Gerrit project (master branch):

```
root@cfg01:~# git add -A

root@cfg01:~# git commit -m 'Switch storage to Gerrit'

root@cfg01:~# git push origin master
```

Step 5 Regenerate the model:

```
root@cfg01:~# salt-call state.sls reclass
```

Step 6 Check the result:

```
root@cfg01:~# cat /srv/salt/reclass/classes/cluster/drivetrain-ha/infra/config/nodes.yml
```

Step 7 Modify the metadata used by your hello formula:

```
root@cfg01:~# cd /root/model_repo/
```

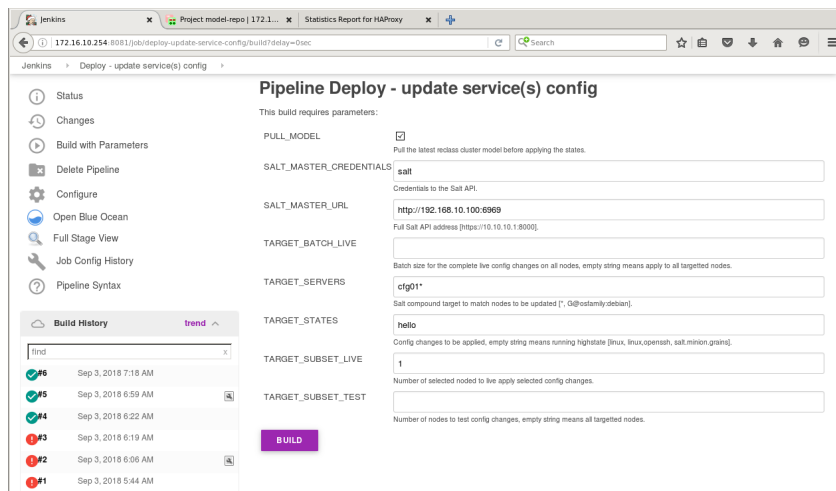
Edit `/root/model_repo/classes/cluster/drivetrain-ha/infra/hello.yml`:

```
parameters:
  hello:
    version: 8
```

Step 8 Add, commit and push to Gerrit project (master branch):

```
root@cfg01:~# git add -A
root@cfg01:~# git commit -m 'Hello ver 8'
root@cfg01:~# git push origin master
```

Step 9 Jenkins GUI - Build with Parameters:



Step 10 Verify the result:

```
root@cfg01:~# cat /etc/hello.conf

This is hello file version 8
```

5.3.9. Use gerrit branching and review GUI

Step 1 Install git hook from Gerrit:

```
root@cfg01:~# cd /root/model_repo/
root@cfg01:~# scp -p -P 29418 admin@172.16.10.104:hooks/commit-msg .git/hooks/
```

Step 2 Modify the metadata used by your hello formula. Edit `/root/model_repo/classes/cluster/drivetrain-ha/infra/hello.yml`:

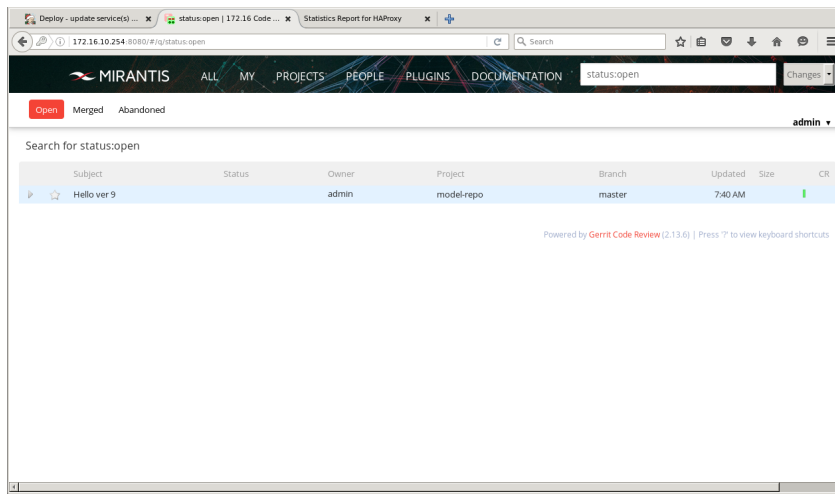
```
parameters:
  hello:
    version: 9
```

Step 3 Add, commit and push to Gerrit project, but this time review branch:

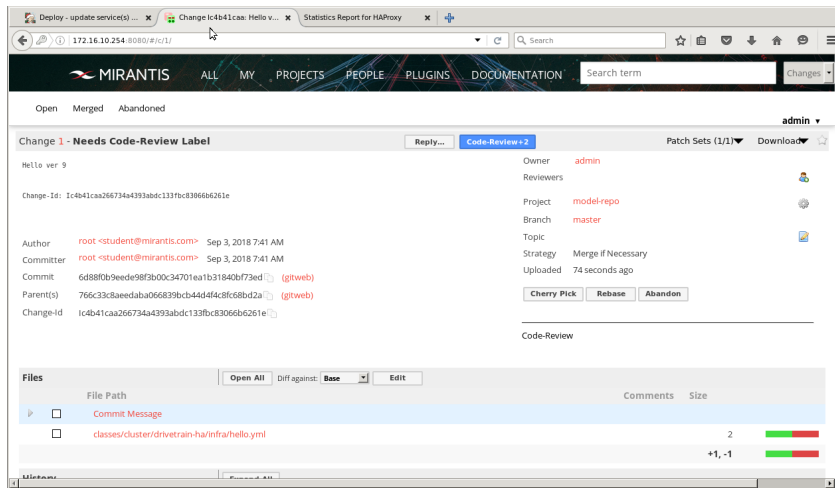
```
root@cfg01:~# git add -A
root@cfg01:~# git commit -m 'Hello ver 9'
root@cfg01:~# git push origin HEAD:refs/for/master
```

Step 4 Go to Gerrit - Review and Submit:

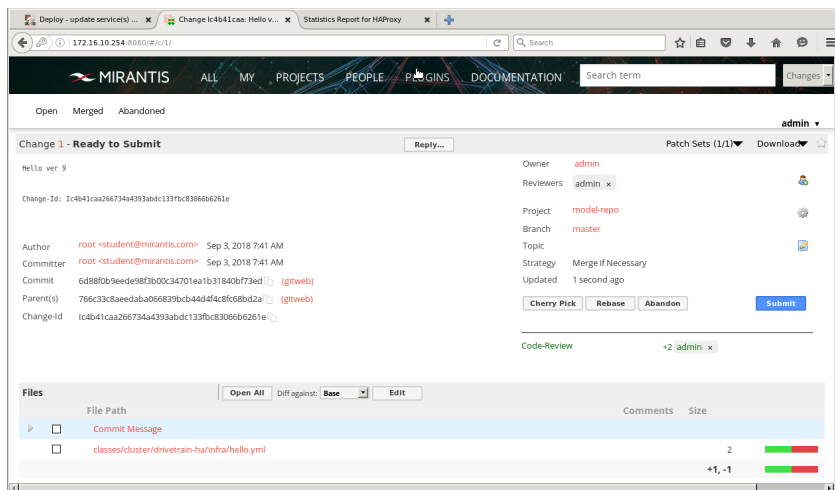
Find your changes in Gerrit:



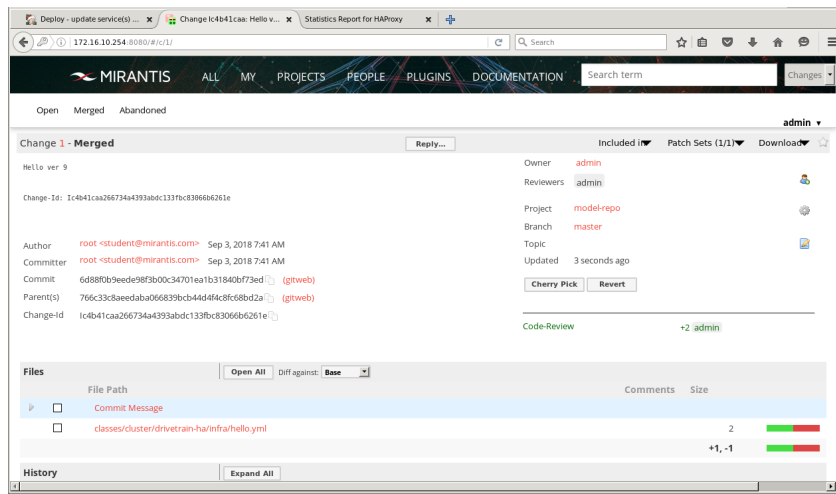
To approve the changes you need to press a “Code-Review +2” button:



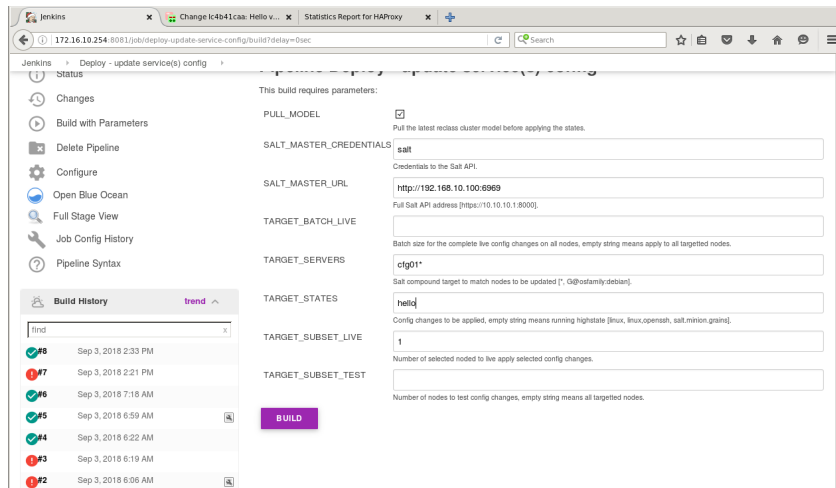
After changes were improved press “Submit” button:



Confirm that changes were merged:



Step 5 Jenkins GUI - Build with Parameters:



Step 6 Verify the result:

```
root@cfg01:~# cat /etc/hello.conf
```

This is hello file version 9