



# MCP Networking

Overlay network solutions for OpenStack and Kubernetes  
supported by Mirantis Cloud Platform (MCP)

[training.mirantis.com](https://training.mirantis.com)

# Objectives

---

- Understand basic networking architectures for MCP clusters (OpenStack, Kubernetes)
- Know the best practices for configuring networks in MCP clusters

# MCP Networking Highlights

---

- OpenStack
  - OpenContrail
  - Neutron with Open vSwitch
- Kubernetes
  - Calico



# OpenStack Networking

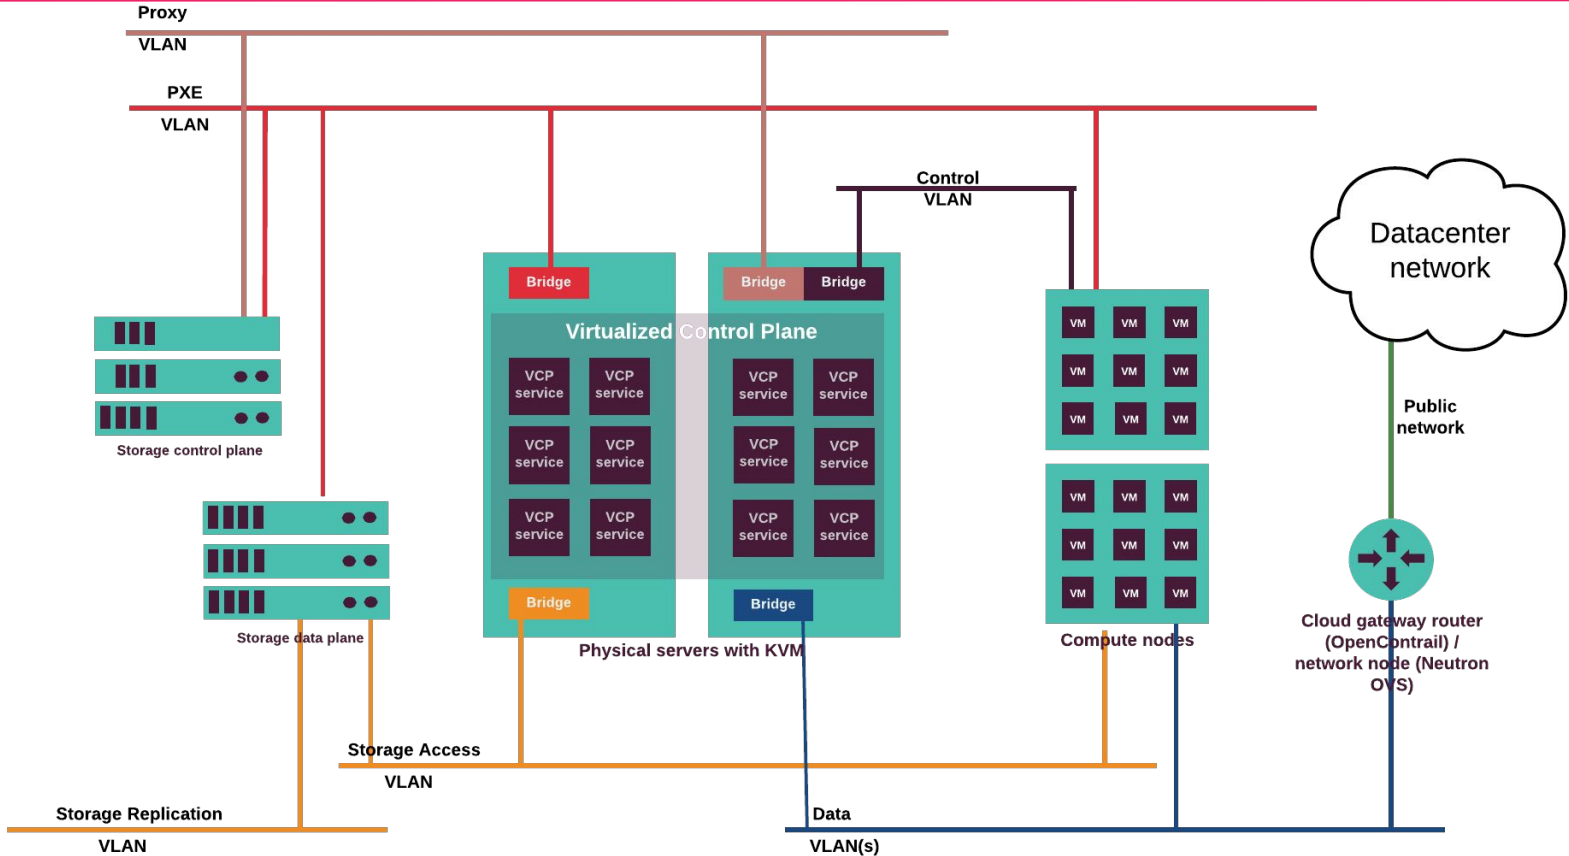
Overlay network solutions for OpenStack supported by MCP

# MCP Network Reference Architecture

---

- OpenContrail (Recommended)
- Neutron with Open vSwitch
  - VxLAN or VLAN
  - DVR or no DVR

# MCP Networking Reference Architecture



# MCP Underlay Networks

---

- PXE / Management
- Public
- Proxy
- Control
- Data
- Storage access (optional)
- Storage replication (optional)

# MCP Networking Recommendations

---

- Isolated virtual or physical networks for PXE, Proxy, and all other traffic.
- Recommended minimal networking configuration:
  - PXE network – one 1 Gbit
  - Other networks – two bonded 10 Gbit
  - Use VLANs to isolate Proxy, Data, and Control networks



# MCP Network Node Configuration

Port	Description	IP Address	VLAN
br-mesh	Tenant overlay traffic (VXLAN)	Routed, Subnet	Leaf switch only
br-mgmt	Openstack and other management traffic	Routed, Subnet	Leaf switch only
br-stor	Storage traffic	Routed, Subnet	Leaf switch only
eth0	PXE Boot traffic	VLAN, Subnet, Default	Global
br-prv	Tenant VLAN traffic bridge	VLAN range	Global
br-floating	External VLAN traffic bridge	VLAN range	Global

# MCP VCP Network Configuration

Port	Description	IP Address	VLAN
br-pxe	PXE Boot traffic	VLAN, Subnet	Global
br-mgmt	Openstack and other management traffic	Routed, Subnet	Leaf switch only
br-stor	Storage traffic	Routed, Subnet	Leaf switch only
eth0	PXE Boot traffic	VLAN, Subnet, Default	Global



# OpenStack with OpenContrail

OpenStack overlay networking with OpenContrail

# OpenContrail only features

---

- Service chaining
- MPLS over UDP/GRE with vMX router
- Multi-site SDN
- Network analytics

# OpenContrail vs Juniper Contrail

---

- Both are based on the same open source code
- Integration with MCP
  - OpenContrail is a part of MCP
  - Extra integration work will need to be done for Juniper Contrail
- Updates
  - Juniper Contrail users are tied to Juniper's release cycle
  - Mirantis can deliver package updates for OpenContrail in a matter of hours

# Clear criteria for using Juniper Contrail

---

- Using Juniper for everything
- OpenStack SDN needs to integrate into the existing hardware and orchestrate special functions, such as
  - extending tenant networks across WAN connectivity with MPLS
  - having those links dynamically change based on events within Contrail

# Mirantis OpenContrail

- There are no 'community' OpenContrail packages
- Mirantis OpenContrail repository
  - <https://github.com/Mirantis/contrail-packages>
- Mirantis consumes upstream OpenContrail repositories as follows:
  - Stores them on a mirror, and mirror output of Mirantis' branches to GitHub for the community
  - Fixes issues and cherry-picks from the master and most recent branches for backports
  - Runs daily builds from stable branches
  - Mirantis' production repositories publicly available via Launchpad

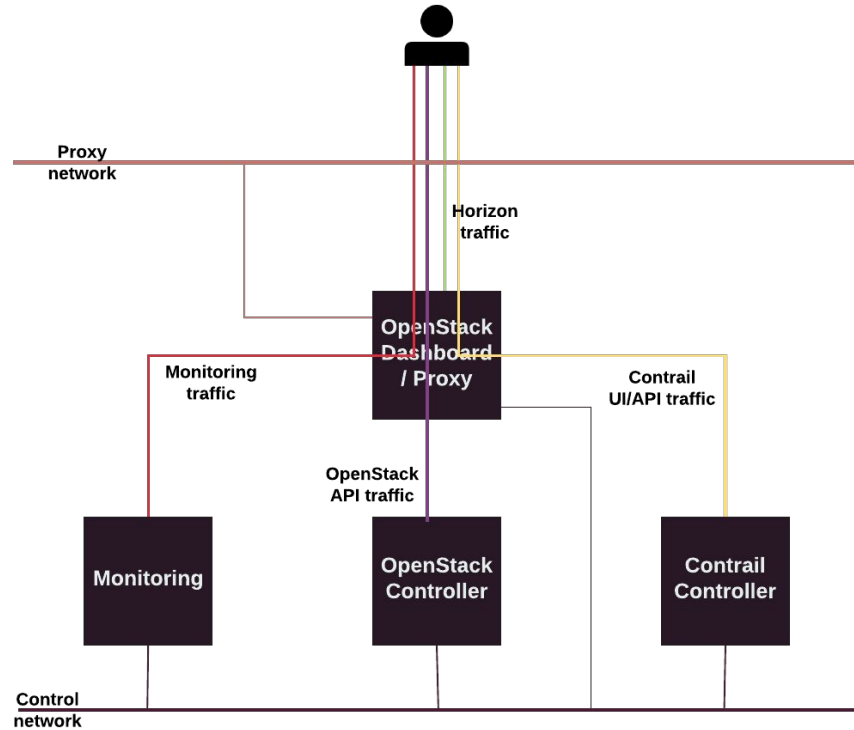
# OpenContrail limitations

---

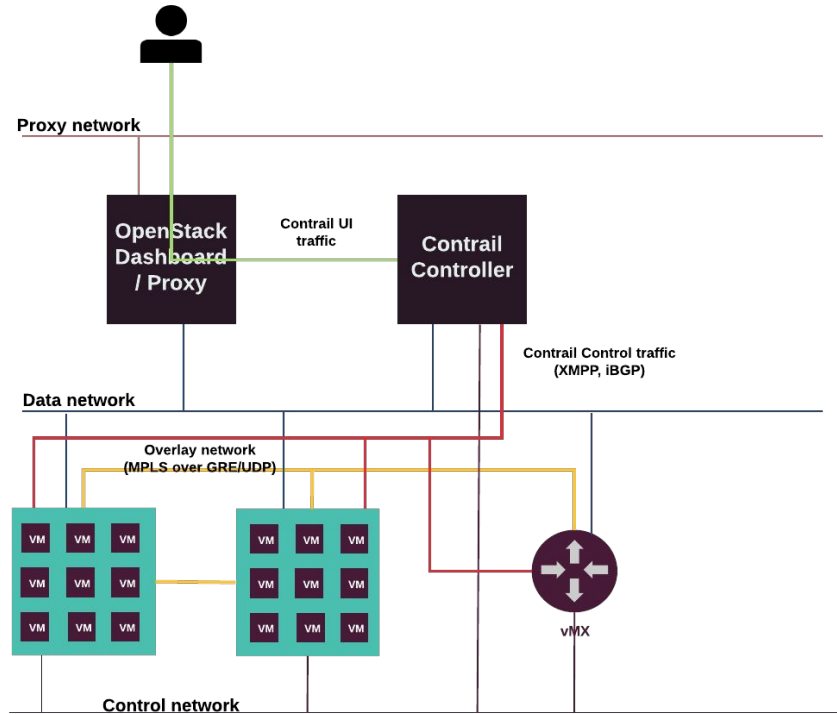
- Exit routers
  - At the moment only Juniper MX and vMX (hardware and software routers) are officially supported



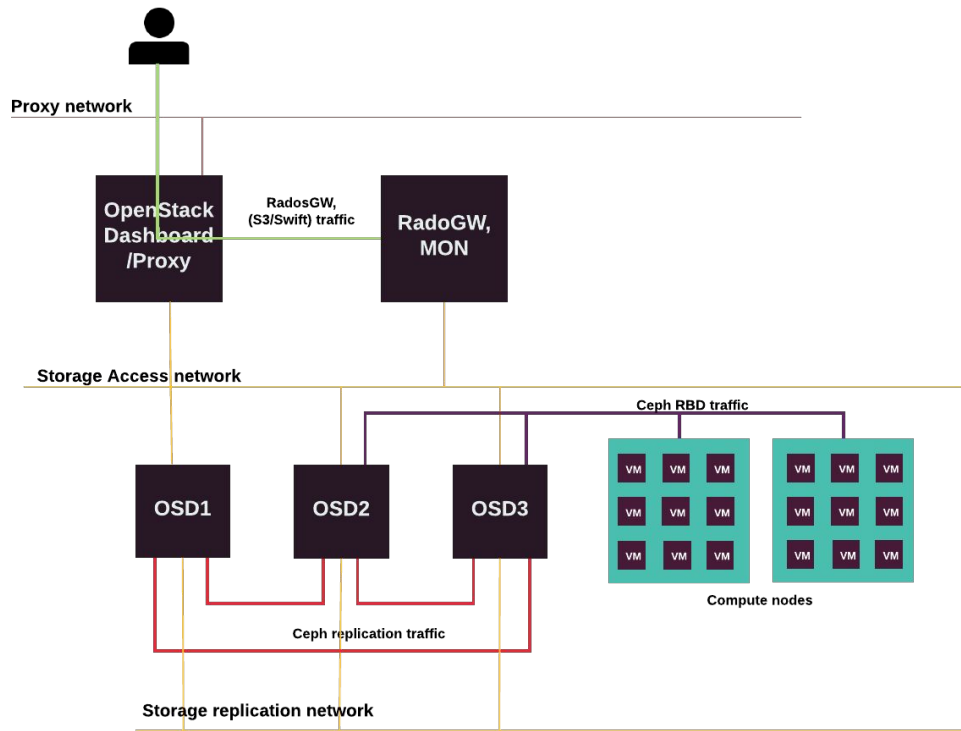
# OpenContrail: User Interface and API traffic



# OpenContrail: SDN traffic



# OpenContrail: Storage traffic





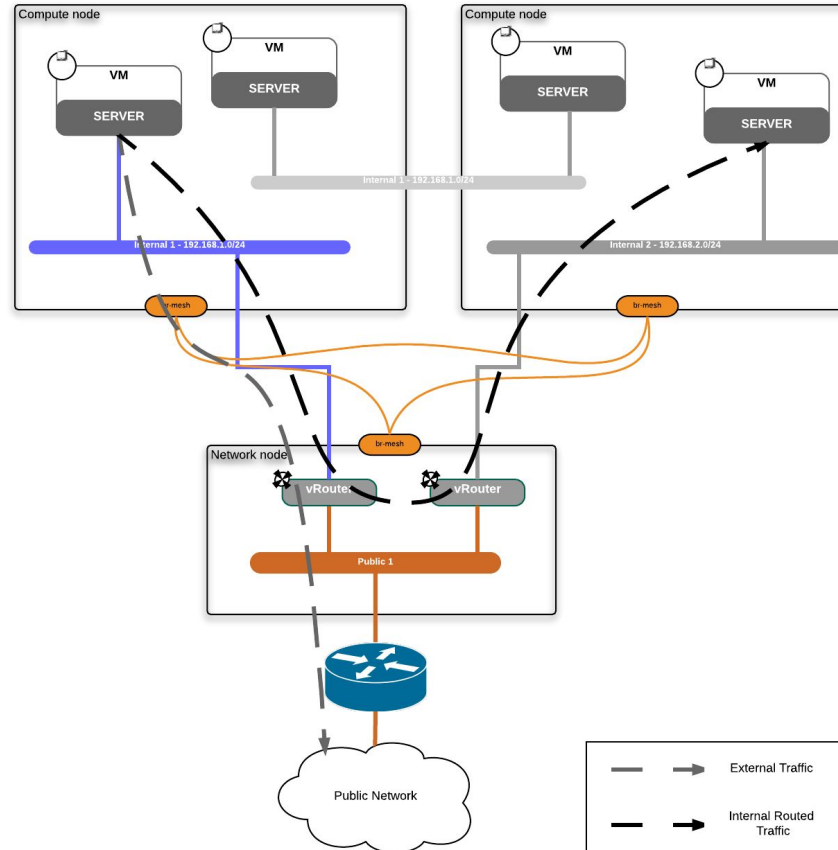
# OpenStack with Open vSwitch


OpenStack overlay networking with Open vSwitch

# MCP Deployment Models for Open vSwitch

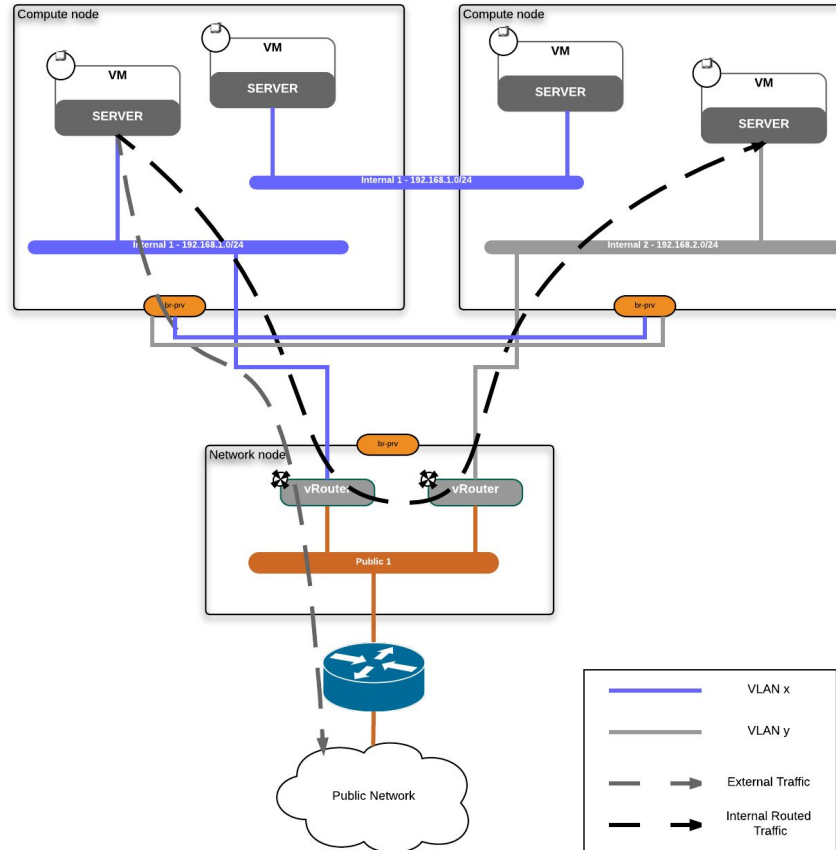
- Neutron VXLAN tenant networks
  - Network Nodes for East-West and North-South
  - DVR for East-West and Network nodes for North-South<sup>\*</sup>
  - DVR for East-West and North-South, Network Node for SNAT<sup>\*\*</sup>
- Neutron VLAN tenant networks
  - Network Nodes for East-West and North-South<sup>\*\*</sup>
  - DVR for East-West and North-South, Network Node for SNAT<sup>\*\*</sup>

# VXLAN, Network Nodes for East-West and North-South



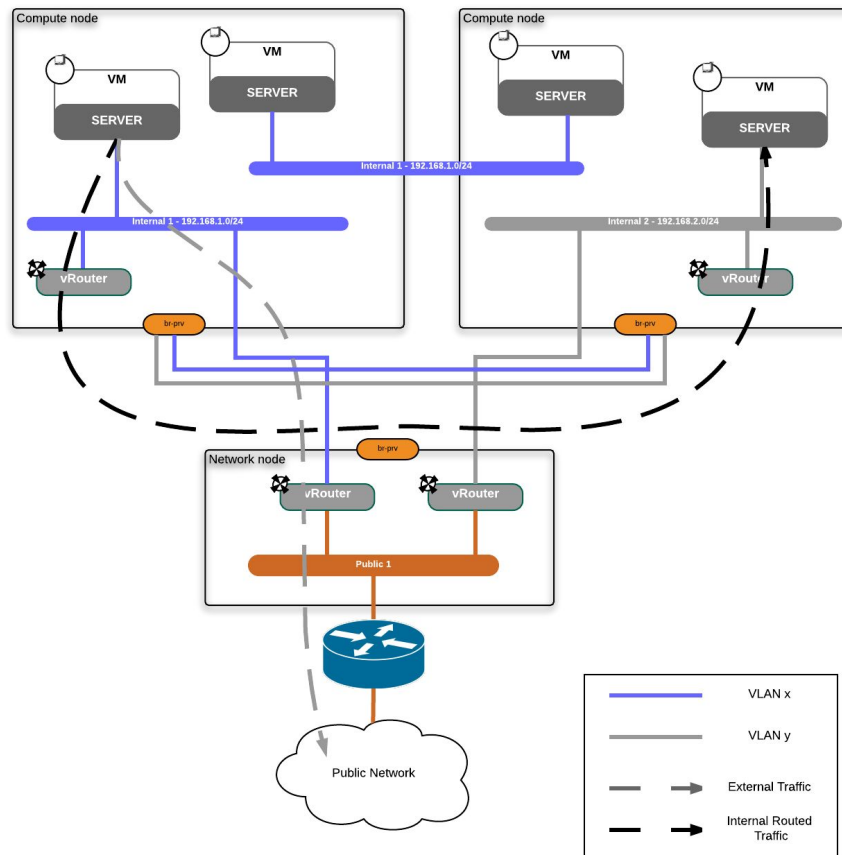


# VLAN, Network Nodes for East-West and North-South





# VLAN, DVR for East-West and North-South, Network Node for SNAT



# MCP Definition of Bridges

```
cat ./system/linux/network/interface/single_ovs_dvr.yml
```

```
parameters:
```

```
  linux:
```

```
    network:
```

```
      bridge: openvswitch
```

```
      interface:
```

```
        primary_interface:
```

```
          name: ${_param:primary_interface}
```

```
          type: eth
```

```
          mtu: ${_param:interface_mtu}
```

```
        tenant_interface:
```

```
          name: ${_param:tenant_interface}
```

```
          type: eth
```

```
          mtu: ${_param:interface_mtu}
```

```
          proto: manual
```

```
br-int:
```

```
  enabled: true
```

```
  mtu: ${_param:interface_mtu}
```

```
  type: ovs_bridge
```

```
br-floating:
```

```
  enabled: true
```

```
  mtu: ${_param:interface_mtu}
```

```
  type: ovs_bridge
```

```
float-to-ex:
```

```
  enabled: true
```

```
  type: ovs_port
```

```
  mtu: 65000
```

```
  bridge: br-floating
```



# Kubernetes with Calico

Kubernetes overlay networking with Calico

# Kubernetes networking backends

---

- Plugins that implement Container Networking Interface (CNI)
- CNI plugins are typically executable scripts + config files
- Which are called by kubelet throughout pod lifecycle management i.e. when pod is created and deleted

# Types of Kubernetes networking

---

- **Underlay**

Pod interfaces are plugged into network fabric (by virtual switch) and traffic is typically bridged across network fabric (works with physical L2 fabric or cloud)

- **Overlay**

Traffic from pod interfaces is sent encapsulated (typically by kernel) over tunnels across network fabric (works with all types of fabrics)

- **Native routing**

Traffic from pods/containers interfaces is natively (by kernel) routed/forwarded from nodes across network fabric

# Native routing vs. other types

---

Native routing has many advantages when compared with underlay and overlay:

- it uses native Linux kernel routing/forwarding capabilities
- no need for additional software switch/router
- it has low overhead
- it works with any type of fabric

# What is Calico

Calico provides:

- SDN – provides connectivity between pods (based on native routing)
  - by exchanging /32 pod IP routes using BGP, and
  - optionally provides IPAM
- Network policy – to restrict traffic between pods (using iptables)

Both functionalities can be used together or separately:

- SDN + policy from Calico
- SDN from Calico alone (without policy)
- SDN from Flannel and policy from Calico

- Control plane
  - BGP routing daemon – bird or gobgp (experimental), it is optional – not needed when Calico is used to provide policy, not SDN
  - Confd – monitor and store config for bird on etcd
- Management plane
  - Felix – main Calico daemon
  - Calicoctl - command line tool to query/modify Calico API primitives/resources on Etcd (supports YAML)
- API (RESTful), internal message bus (REST pub-sub) and persistent storage
  - Etcd – used as external API server, internal message bus and persistent storage, all the internal and external communication on Calico goes through etcd



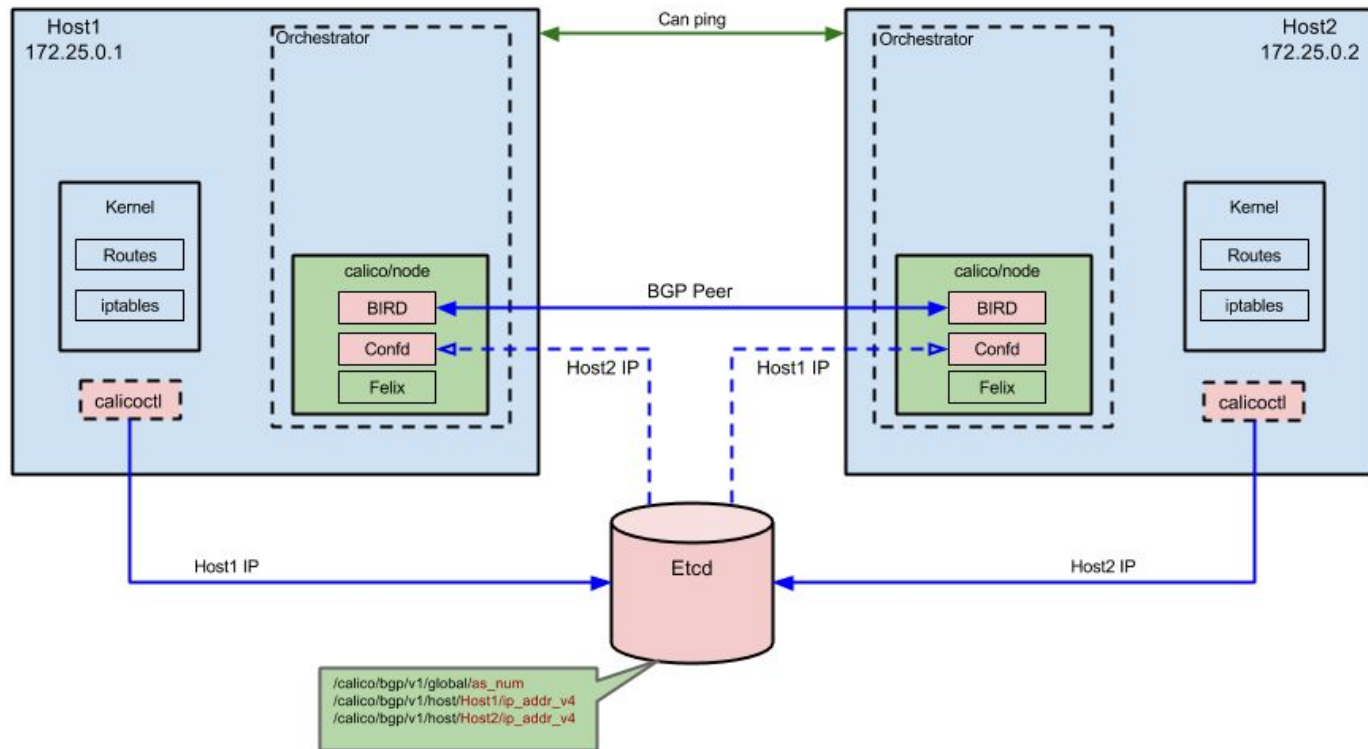
# Calico architecture - Data plane

---


- standard kernel forwarding (native L3 routing)
- iptables – policy
- veth is set with arp proxy enabled

# Calico integration with Kubernetes

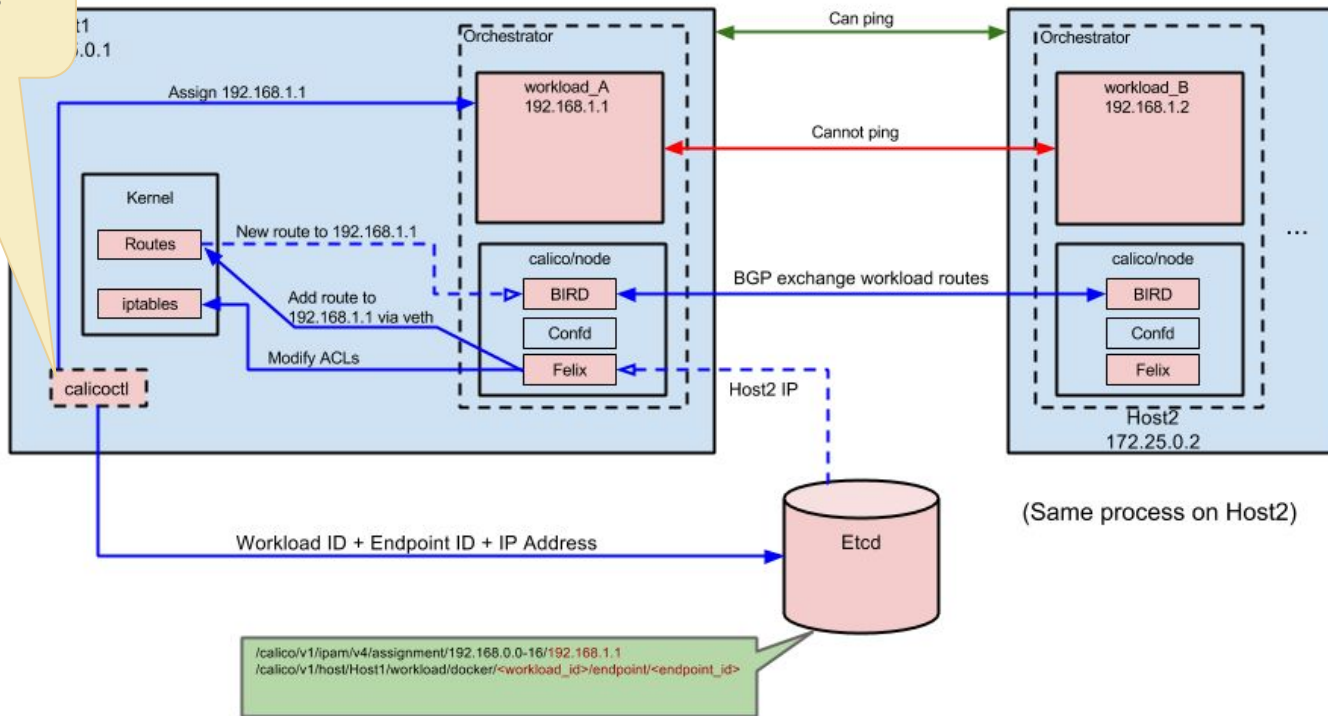
- SDN
  - calico CNI (SDN) plugin
  - calico CNI IPAM plugin (optional)
  - CNI config file
- Policy
  - policy controller translates Kubernetes policy (from Kubernetes API) into Calico policy stored on etcd
- Service IPs and local internal load balancing
  - are configured by Kubernetes on its own, using kube-proxy and iptables

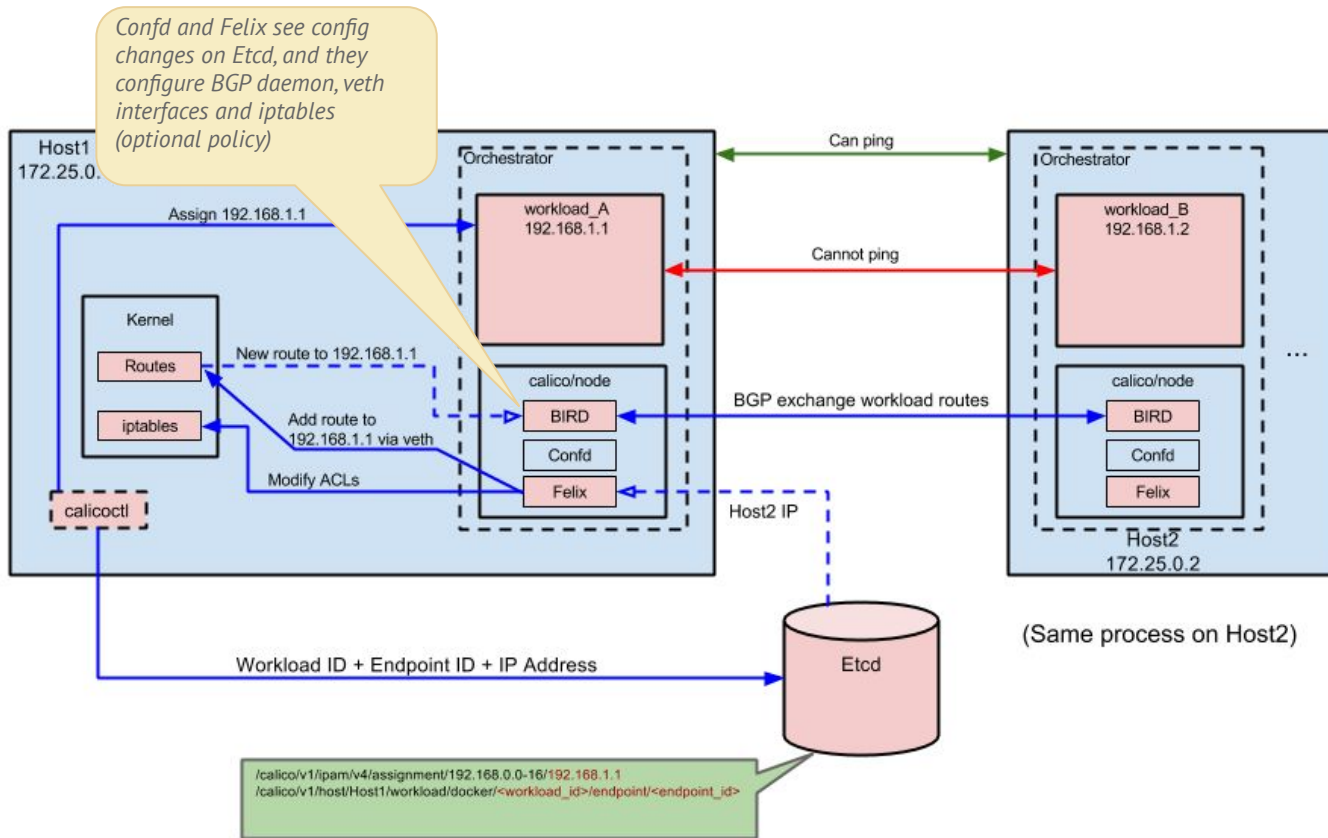


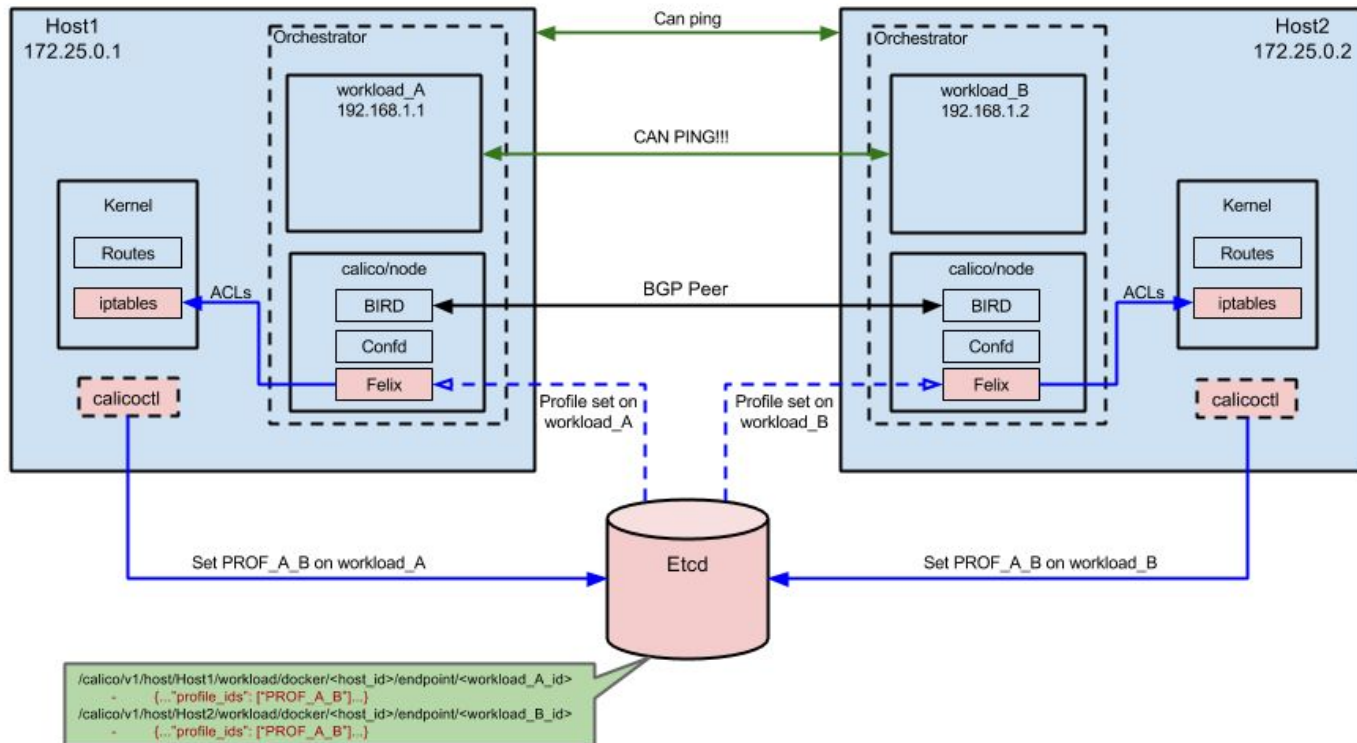
*Pods (workload A and B) are created by Kubernetes - orchestrator).  
Pause Docker container is connected to Docker network none*



Kubernetes (kubelet) calls Calico CNI plugins. IPAM plugin is optional. Plugins call Calico API (etcd) to create configuration.







# Calico deployment with MCP

- Docker container 'calico/node' runs on every Kubernetes node
  - it contains Calico control and management plane daemons
  - it is started using host network (node IP), so there is no egg chicken problem
- Since kubelet runs as a container, directory with CNI plugins and config file must be mounted as volume on kubelet container so they can be called/executed by kubelet
- Salt formula for Kubernetes does the magic
  - <https://github.com/salt-formulas/salt-formula-kubernetes>